

--	--	--	--	--	--	--	--	--	--

Day: ☐ Tuesday ☐ Wednesday ☐ FridayTime: ☐ 10:00 - 10:50 ☐ 11:00 - 11:50 ☐ 16:00 - 16:50 ☐ 18:00 - 18:50

Condition Variables

Introduction

Topic to be covered in this tutorial include:

- Explore some real code that uses locks and condition variables to implement various forms of the producer/consumer queue.

Acknowledgement

This tutorial was adapted from OSTEP book written by Remzi and Andrea Arpaci-Dusseau at the University of Wisconsin. This free OS book is available at <http://www.ostep.org>.

Getting Started

1. Logging in to the Linux server

- Start the SSH client, e.g., MobaXterm or Xshell.
- Log in to the Linux server using the following details:

Host Name: gateway.cs.cityu.edu.hk

User Name: your EID (e.g., cctom2)

Password: your password

🔒 Your password will not be shown on the screen as you type it, not even as a row of stars (*****).

NOTE: The shell will always give you a prompt if it is ready to accept commands. The shell prompt normally ends in a \$ sign as we use in this tutorial. Some shell prompts end in % or > instead. Never copy/type the shell prompt used in this tutorial.

NOTE: Please don't forget to log out (use the `exit` command) after you finish your work.

2. Copying the Example Code

This tutorial lets you explore some real code that uses locks and condition variables to implement various forms of the producer/consumer queue, which are in the directory `/public/cs3103/tutorial5`.

Start by copying them to a directory in which you plan to do your work. For example, to copy `tutorial5` directory to your current directory and change to it, enter:

```
$ cp -rf /public/cs3103/tutorial5 .  
$ cd tutorial5
```

The last dot/period (.) indicates the current directory as the destination.

Introduction

The first step is to download the code and type `make` to build all the variants. You should see three:

- `main-one-cv-while.c`: The producer/consumer problem solved with a single condition variable.
- `main-two-cvs-if.c`: Same but with two condition variables and using an `if` to check whether to sleep.
- `main-two-cvs-while.c`: Same but with two condition variables and `while` to check whether to sleep. This is the correct version.

It's also useful to look at `pc-header.h`, which contains common code for all of these different main programs, and the `Makefile` so as to build the code properly.

Each program takes the following flags:

```
-l <number of items each producer produces> (lower-case letter L)
-m <size of the shared producer/consumer buffer>
-p <number of producers>
-c <number of consumers>
-P <sleep string: how producer should sleep at various points>
-C <sleep string: how consumer should sleep at various points>
-v [verbose flag: trace what is happening and print it]
-t [timing flag: time entire execution and print total time]
```

The first four arguments are relatively self-explanatory: `-l` specifies how many times each producer should loop (and thus how many data items each producer produces), `-m` controls the size of the shared buffer (greater than or equal to one), and `-p` and `-c` set how many producers and consumers there are, respectively.

What is more interesting are the two sleep strings, one for producers, and one for consumers. These flags allow you to make each thread sleep at certain points in an execution and thus switch to other threads; doing so allows you to play with each solution and perhaps pinpoint specific problems or study other facets of the producer/consumer problem.

The string is specified as follows. If there are three producers, for example, the string should specify sleep times for each producer separately, with a colon separator. The sleep string for these three producers would look something like this:

```
sleep_string_for_p0:sleep_string_for_p1:sleep_string_for_p2
```

Each sleep string, in turn, is a comma-separated list, which is used to decide how much to sleep at each sleep point in the code. To understand this per-producer or per-consumer sleep string better, let's look at the code in

`main-two-cvs-while.c`, specifically at the producer code. In this code snippet, a producer thread loops for a while, putting elements into the shared buffer via calls to `do_fill()`. Around this filling routine are some waiting and signaling code to ensure the buffer is not full when the producer tries to fill it.

```

void *producer(void *arg) {
    int id = (int) arg;
    // make sure each producer produces unique values
    int base = id * loops;
    int i;
    for (i = 0; i < loops; i++) {
        p0;
        Mutex_lock(&m);          p1;
        while (num_full == max) { p2;
            Cond_wait(&empty, &m); p3;
        }
        do_fill(base + i);        p4;
        Cond_signal(&fill);       p5;
        Mutex_unlock(&m);         p6;
    }
    return NULL;
}

```

As you can see from the code, there are a number of points labeled p0, p1, etc. These points are where the code can be put to sleep. The consumer code (not shown here) has similar wait points (c0, etc.).

Specifying a sleep string for a producer is easy: just identify how long the producer should sleep at each point p0, p1, ..., p6. For example, the string 1,0,0,0,0,0,0 specifies that the producer should sleep for 1 second at marker p0 (before grabbing the lock), and then not at all each time through the loop.

Now let's show the output of running one of these programs. To begin, let's assume that the user runs with just one producer and one consumer. Let's not add any sleeping at all (this is the default behavior). The buffer size, in this example, is set to 2 (-m 2).

First, let's build the code:

```
$ make
```

Now we can run something:

```
$ ./main-two-cvs-while -l 3 -m 2 -p 1 -c 1 -v
```

In this case, if you trace the code (with the verbose flag, -v), you will get the following output (or something like it) on the screen:

```

NF          P0 C0
0 [ *---   --- ] p0
0 [ *---   --- ]   c0
0 [ *---   --- ] p1
1 [ u  0 f--- ] p4

```

```

1 [u  0 f--- ] p5
1 [u  0 f--- ] p6
1 [u  0 f--- ] p0
1 [u  0 f--- ] c1
0 [  --- *--- ] c4
0 [  --- *--- ] c5
0 [  --- *--- ] c6
0 [  --- *--- ] c0
0 [  --- *--- ] p1
1 [f--- u  1 ] p4
1 [f--- u  1 ] p5
1 [f--- u  1 ] p6
1 [f--- u  1 ] p0
1 [f--- u  1 ] c1
0 [*---  --- ] c4
0 [*---  --- ] c5
0 [*---  --- ] c6
0 [*---  --- ] c0
0 [*---  --- ] p1
1 [u  2 f--- ] p4
1 [u  2 f--- ] p5
1 [u  2 f--- ] p6
1 [u  2 f--- ] c1
0 [  --- *--- ] c4
0 [  --- *--- ] c5
0 [  --- *--- ] c6
1 [f--- uEOS ] [main: added end-of-stream marker]
1 [f--- uEOS ] c0
1 [f--- uEOS ] c1
0 [*---  --- ] c4
0 [*---  --- ] c5
0 [*---  --- ] c6
Consumer consumption:
C0 -> 3

```

Before describing what is happening in this simple example, let's understand the depiction of the shared buffer in the output, as is shown on the left. At first it is empty (an empty slot indicated by ---, and the two empty slots shown as [*--- ---]); the output also shows the number of entries in the buffer (num_full), which starts at 0. Then, after P0 puts a value in it, its state changes to [u 0 f---] (and num_full changes to 1). You might also notice a few additional markers here; the u marker shows where the use_ptr is (this is where the next consumer to consume a value will get something from); similarly, the f marker shows where the fill_ptr is

(this is where the next producer will produce a value). When you see the * marker, it just means the use_ptr and fill_ptr are pointing to the same slot.

Along the right you can see the trace of which step each producer and consumer is about to execute. For example, the producer grabs the lock (p1), and then, because the buffer has an empty slot, produces a value into it (p4). It then continues until the point where it releases the lock (p6) and then tries to reacquire it. In this example, the consumer acquires the lock instead and consumes the value (c1, c4). Study the trace further to understand how the producer/consumer solution works with a single producer and consumer.

Now let's add some pauses to change the behavior of the trace. In this case, let's say we want to make the producer sleep so that the consumer can run first. We can accomplish this as follows:

```
$ ./main-two-cvs-while -l 1 -m 2 -p 1 -c 1 -P 1,0,0,0,0,0,0 -C 0 -v
```

The results:

```

NF          P0 C0
0 [*--- --- ] p0
0 [*--- --- ]   c0
0 [*--- --- ]   c1
0 [*--- --- ]   c2
0 [*--- --- ] p1
1 [u  0 f--- ] p4
1 [u  0 f--- ] p5
1 [u  0 f--- ] p6
1 [u  0 f--- ]   c3
0 [ --- *--- ]   c4
0 [ --- *--- ]   c5
0 [ --- *--- ]   c6
0 [f--- uEOS ]   c0
1 [f--- uEOS ] [main: added end-of-stream marker]
1 [f--- uEOS ]   c1
0 [*--- --- ]   c4
0 [*--- --- ]   c5
0 [*--- --- ]   c6

```

Consumer consumption:

```
C0 -> 1
```

As you can see, the producer hits the p0 marker in the code and then grabs the first value from its sleep specification, which in this case is 1, and thus each sleeps for 1 second before even trying to grab the lock. Thus, the consumer gets to run, grabs the lock, but finds the queue empty, and thus sleeps (releasing the lock). The producer then runs (eventually), and all proceeds as you might expect.

Do note: a sleep specification must be given for each producer and consumer. Thus, if you create two producers and three consumers (with `-p 2 -c 3`, you must specify sleep strings for each (e.g., `-P 0:1` or `-C 0,1,2:0:3,3,3,1,1,1`). Sleep strings can be shorter than the number of sleep points in the code; the remaining sleep slots are initialized to be zero.

Questions

See the answer sheet file. All questions should be answered on the separate answer sheet provided.